

```

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

# Importing California Housing Prices Dataset

from pathlib import Path
import tarfile
import urllib.request

def load_housing_data():
    tarball_path = Path("datasets/housing.tgz")
    if not tarball_path.is_file():
        Path("datasets").mkdir(parents=True, exist_ok=True)
        url = "https://github.com/ageron/data/raw/main/housing.tgz"
        urllib.request.urlretrieve(url, tarball_path)
        with tarfile.open(tarball_path) as housing_tarball:
            housing_tarball.extractall(path="datasets")
    return pd.read_csv(Path("datasets/housing/housing.csv"))

```

```
housing = load_housing_data()
```

```
housing
```

	longitude	latitude	housing_median_age	total_rooms
total_bedrooms \				
0	-122.23	37.88	41.0	880.0
129.0				
1	-122.22	37.86	21.0	7099.0
1106.0				
2	-122.24	37.85	52.0	1467.0
190.0				
3	-122.25	37.85	52.0	1274.0
235.0				
4	-122.25	37.85	52.0	1627.0
280.0				
...	...	...	...	...
...				
20635	-121.09	39.48	25.0	1665.0
374.0				
20636	-121.21	39.49	18.0	697.0
150.0				
20637	-121.22	39.43	17.0	2254.0
485.0				
20638	-121.32	39.43	18.0	1860.0
409.0				
20639	-121.24	39.37	16.0	2785.0
616.0				

	population	households	median_income	median_house_value \
0	322.0	126.0	8.3252	452600.0
1	2401.0	1138.0	8.3014	358500.0
2	496.0	177.0	7.2574	352100.0
3	558.0	219.0	5.6431	341300.0
4	565.0	259.0	3.8462	342200.0
...	...	...	...	...
20635	845.0	330.0	1.5603	78100.0
20636	356.0	114.0	2.5568	77100.0
20637	1007.0	433.0	1.7000	92300.0
20638	741.0	349.0	1.8672	84700.0
20639	1387.0	530.0	2.3886	89400.0

	ocean_proximity
0	NEAR BAY
1	NEAR BAY
2	NEAR BAY
3	NEAR BAY
4	NEAR BAY
...	...
20635	INLAND
20636	INLAND
20637	INLAND
20638	INLAND
20639	INLAND

[20640 rows x 10 columns]

Data

Preprocessing

df=housing.copy()

housing.head()

	longitude	latitude	housing_median_age	total_rooms
total_bedrooms \				
0	-122.23	37.88	41.0	880.0
129.0				
1	-122.22	37.86	21.0	7099.0
1106.0				
2	-122.24	37.85	52.0	1467.0
190.0				
3	-122.25	37.85	52.0	1274.0
235.0				
4	-122.25	37.85	52.0	1627.0
280.0				

	population	households	median_income	median_house_value
ocean_proximity				

0	322.0	126.0	8.3252	452600.0
NEAR BAY				
1	2401.0	1138.0	8.3014	358500.0
NEAR BAY				
2	496.0	177.0	7.2574	352100.0
NEAR BAY				
3	558.0	219.0	5.6431	341300.0
NEAR BAY				
4	565.0	259.0	3.8462	342200.0
NEAR BAY				

```
housing.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude             20640 non-null  float64
1   latitude              20640 non-null  float64
2   housing_median_age    20640 non-null  float64
3   total_rooms           20640 non-null  float64
4   total_bedrooms        20433 non-null  float64
5   population            20640 non-null  float64
6   households            20640 non-null  float64
7   median_income         20640 non-null  float64
8   median_house_value    20640 non-null  float64
9   ocean_proximity       20640 non-null  object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
```

```
housing.isnull().sum()
```

longitude	0
latitude	0
housing_median_age	0
total_rooms	0
total_bedrooms	207
population	0
households	0
median_income	0
median_house_value	0
ocean_proximity	0
dtype:	int64

total\_bedrooms contains 207 missing value we have to impute them before training machine learning model

```
housing.describe()
```

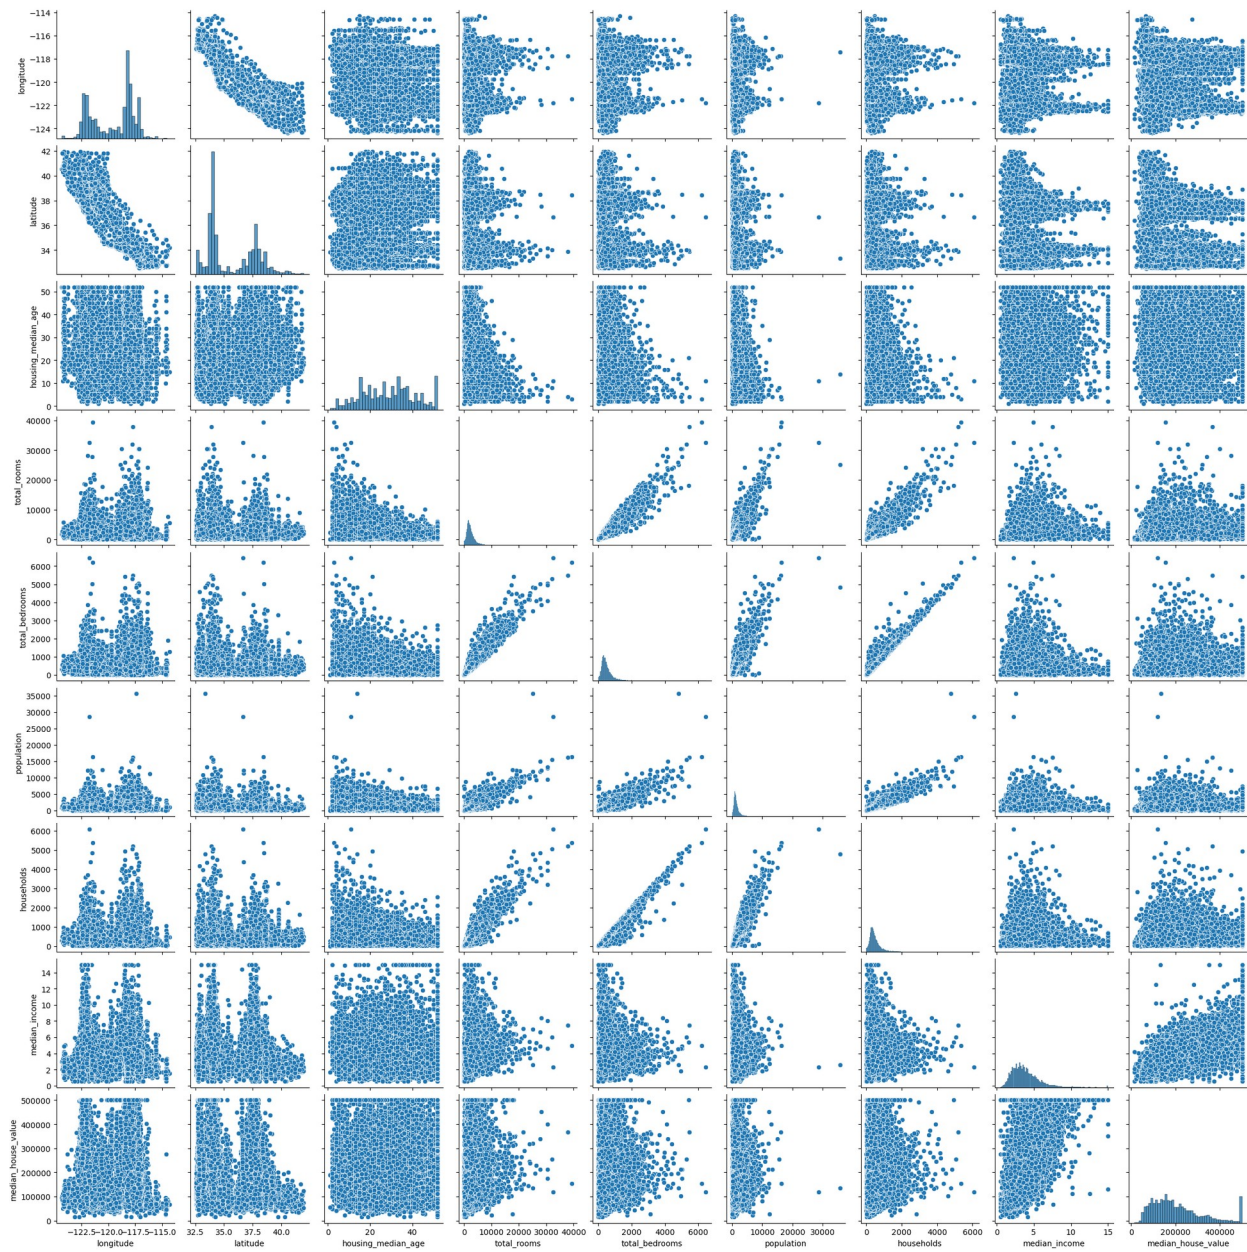
	longitude	latitude	housing_median_age	total_rooms	\
count	20640.000000	20640.000000	20640.000000	20640.000000	
mean	-119.569704	35.631861	28.639486	2635.763081	
std	2.003532	2.135952	12.585558	2181.615252	
min	-124.350000	32.540000	1.000000	2.000000	
25%	-121.800000	33.930000	18.000000	1447.750000	
50%	-118.490000	34.260000	29.000000	2127.000000	
75%	-118.010000	37.710000	37.000000	3148.000000	
max	-114.310000	41.950000	52.000000	39320.000000	

	total_bedrooms	population	households	median_income	\
count	20433.000000	20640.000000	20640.000000	20640.000000	
mean	537.870553	1425.476744	499.539680	3.870671	
std	421.385070	1132.462122	382.329753	1.899822	
min	1.000000	3.000000	1.000000	0.499900	
25%	296.000000	787.000000	280.000000	2.563400	
50%	435.000000	1166.000000	409.000000	3.534800	
75%	647.000000	1725.000000	605.000000	4.743250	
max	6445.000000	35682.000000	6082.000000	15.000100	

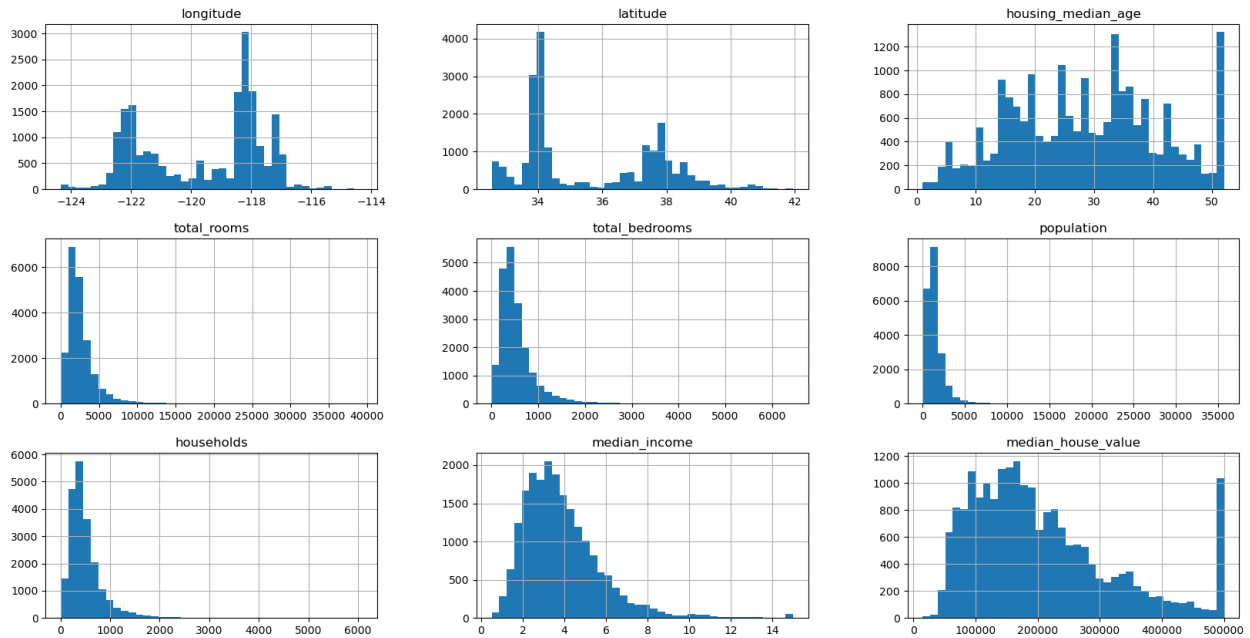
	median_house_value
count	20640.000000
mean	206855.816909
std	115395.615874
min	14999.000000
25%	119600.000000
50%	179700.000000
75%	264725.000000
max	500001.000000

```
sns.pairplot(housing)
```

```
<seaborn.axisgrid.PairGrid at 0x14e40eed0>
```



```
%matplotlib inline
housing.hist(bins=40,figsize=(20,10))
plt.show()
```




total\_rooms, total\_bedrooms and household are right tailed data to make them normal distributed we can apply natural log transformation or log base 10 transformation on them

median\_house\_age has almost normal distribution and normal\_house\_value is little bit right tailed where we can apply natural log to make it normal distributed

```
# let's see how actual california dataset looks like
sns.scatterplot(x=housing.longitude,y=housing.latitude,hue=housing['ocean_proximity'])
plt.title('California Housing Price Map')
plt.show()
```





```
# checking the correlation between different features
corr=housing.drop(columns=['ocean_proximity']).corr()
corr
```

	longitude	latitude	housing_median_age
total_rooms \			
longitude	1.000000	-0.924664	-0.108197
0.044568			
latitude	-0.924664	1.000000	0.011173
0.036100			-
housing_median_age	-0.108197	0.011173	1.000000
0.361262			-
total_rooms	0.044568	-0.036100	-0.361262
1.000000			
total_bedrooms	0.069608	-0.066983	-0.320451
0.930380			
population	0.099773	-0.108785	-0.296244
0.857126			
households	0.055310	-0.071035	-0.302916
0.918484			
median_income	-0.015176	-0.079809	-0.119034
0.198050			
median_house_value	-0.045967	-0.144160	0.105623


0.134153

	total_bedrooms	population	households	
median_income \				
longitude	0.069608	0.099773	0.055310	-
0.015176				
latitude	-0.066983	-0.108785	-0.071035	-
0.079809				
housing_median_age	-0.320451	-0.296244	-0.302916	-
0.119034				
total_rooms	0.930380	0.857126	0.918484	
0.198050				
total_bedrooms	1.000000	0.877747	0.979728	-
0.007723				
population	0.877747	1.000000	0.907222	
0.004834				
households	0.979728	0.907222	1.000000	
0.013033				
median_income	-0.007723	0.004834	0.013033	
1.000000				
median_house_value	0.049686	-0.024650	0.065843	
0.688075				

	median_house_value
longitude	-0.045967
latitude	-0.144160
housing_median_age	0.105623
total_rooms	0.134153
total_bedrooms	0.049686
population	-0.024650
households	0.065843
median_income	0.688075
median_house_value	1.000000

```
plt.figure(figsize=(20,8))  
sns.heatmap(corr,annot=True)  
plt.show()
```





From above visualization we can conclude that some features are highly correlated with each other which cause multicollinearity and as we know some machine learning algorithms doesn't work well with multicollinearity.

*# To get rid of multicollinearity we try some new features in our dataset*

```
housing["rooms_per_household"] =
housing["total_rooms"]/housing["households"]
housing["bedrooms_per_room"] =
housing["total_bedrooms"]/housing["total_rooms"]
housing["population_per_household"]=housing["population"]/housing["households"]
```

```
corr1=housing.drop(columns=['ocean_proximity']).corr()
corr1
```

	longitude	latitude	housing_median_age	\
longitude	1.000000	-0.924664	-0.108197	
latitude	-0.924664	1.000000	0.011173	
housing_median_age	-0.108197	0.011173	1.000000	
total_rooms	0.044568	-0.036100	-0.361262	
total_bedrooms	0.069608	-0.066983	-0.320451	
population	0.099773	-0.108785	-0.296244	
households	0.055310	-0.071035	-0.302916	
median_income	-0.015176	-0.079809	-0.119034	
median_house_value	-0.045967	-0.144160	0.105623	
rooms_per_household	-0.027540	0.106389	-0.153277	
bedrooms_per_room	0.092657	-0.113815	0.136089	
population_per_household	0.002476	0.002366	0.013191	

	total_rooms	total_bedrooms	population
total_rooms	1.000000	0.930000	0.860000
total_bedrooms	0.930000	1.000000	0.880000
population	0.860000	0.880000	1.000000

households \				
longitude	0.044568	0.069608	0.099773	
0.055310				
latitude	-0.036100	-0.066983	-0.108785	-
0.071035				
housing_median_age	-0.361262	-0.320451	-0.296244	-
0.302916				
total_rooms	1.000000	0.930380	0.857126	
0.918484				
total_bedrooms	0.930380	1.000000	0.877747	
0.979728				
population	0.857126	0.877747	1.000000	
0.907222				
households	0.918484	0.979728	0.907222	
1.000000				
median_income	0.198050	-0.007723	0.004834	
0.013033				
median_house_value	0.134153	0.049686	-0.024650	
0.065843				
rooms_per_household	0.133798	0.001538	-0.072213	-
0.080598				
bedrooms_per_room	-0.187900	0.084238	0.035319	
0.065087				
population_per_household	-0.024581	-0.028355	0.069863	-
0.027309				


	median_income	median_house_value \
longitude	-0.015176	-0.045967
latitude	-0.079809	-0.144160
housing_median_age	-0.119034	0.105623
total_rooms	0.198050	0.134153
total_bedrooms	-0.007723	0.049686
population	0.004834	-0.024650
households	0.013033	0.065843
median_income	1.000000	0.688075
median_house_value	0.688075	1.000000
rooms_per_household	0.326895	0.151948
bedrooms_per_room	-0.615661	-0.255880
population_per_household	0.018766	-0.023737

	rooms_per_household	bedrooms_per_room \
longitude	-0.027540	0.092657
latitude	0.106389	-0.113815
housing_median_age	-0.153277	0.136089
total_rooms	0.133798	-0.187900
total_bedrooms	0.001538	0.084238
population	-0.072213	0.035319
households	-0.080598	0.065087
median_income	0.326895	-0.615661

median_house_value	0.151948	-0.255880
rooms_per_household	1.000000	-0.416952
bedrooms_per_room	-0.416952	1.000000
population_per_household	-0.004852	0.002938

	population_per_household
longitude	0.002476
latitude	0.002366
housing_median_age	0.013191
total_rooms	-0.024581
total_bedrooms	-0.028355
population	0.069863
households	-0.027309
median_income	0.018766
median_house_value	-0.023737
rooms_per_household	-0.004852
bedrooms_per_room	0.002938
population_per_household	1.000000

```
plt.figure(figsize=(20,8))
sns.heatmap(corr1,annot=True)
plt.show()
```



Now we can drop these columns from housing and Let's go to next step where we make our dataset fully ready to feed into machine learning algorithm

```
housing.drop(columns=['total_rooms', 'total_bedrooms', 'population', 'households'], inplace=True)
housing.head()
```

	longitude	latitude	housing_median_age	median_income
0	-122.23	37.88	41.0	8.3252
1	-122.22	37.86	21.0	8.3014
2	-122.24	37.85	52.0	7.2574
3	-122.25	37.85	52.0	5.6431
4	-122.25	37.85	52.0	3.8462

	ocean_proximity	rooms_per_household	bedrooms_per_room
0	NEAR BAY	6.984127	0.146591
1	NEAR BAY	6.238137	0.155797
2	NEAR BAY	8.288136	0.129516
3	NEAR BAY	5.817352	0.184458
4	NEAR BAY	6.281853	0.172096

	population_per_household
0	2.555556
1	2.109842
2	2.802260
3	2.547945
4	2.181467

housing[housing['median\_house\_value']>=500000]

	longitude	latitude	housing_median_age	median_income
89	-122.27	37.80	52.0	1.2434
103	-118.47	33.99	24.0	2.9750
105	-118.50	33.97	29.0	5.1280
107	-118.39	34.08	27.0	3.8088
132	-122.34	37.55	44.0	6.9533
...	...	...	...	...
20494	-118.12	34.13	52.0	11.7060
20500	-121.93	37.66	24.0	8.3337
20511	-122.05	37.31	25.0	9.2298
20515	-117.36	33.17	24.0	2.3182
20535	-122.43	37.80	52.0	4.5399

	median_house_value	ocean_proximity	rooms_per_household
89	500001.0	NEAR BAY	2.929412
103	500001.0	<1H OCEAN	3.456731
105	500001.0	<1H OCEAN	3.932471
107	500001.0	<1H OCEAN	4.345395
132	500001.0	NEAR OCEAN	7.608025
...	...	...	...
20494	500001.0	<1H OCEAN	8.975535

20500	500001.0	<1H OCEAN	7.915000
20511	500001.0	<1H OCEAN	7.237676
20515	500001.0	NEAR OCEAN	5.574932
20535	500001.0	NEAR BAY	4.898601


	bedrooms_per_room	population_per_household
89	0.313253	4.658824
103	0.315716	1.598558
105	0.295214	1.662356
107	0.258895	1.753289
132	0.133063	2.601852
...	...	...
20494	0.116184	2.981651
20500	0.133923	2.702500
20511	0.130868	2.790493
20515	0.216031	2.212534
20535	0.221984	1.667832

[992 rows x 9 columns]

```
capped_values = housing[housing['median_house_value'] >= 500001]
print(f"Capped values: {len(capped_values)} out of {len(housing)}
({len(capped_values)/len(housing)*100:.2f}%)")
```

```
plt.figure(figsize=(8, 6))
plt.hist(housing['median_house_value'], bins=40, edgecolor='black')
plt.axvline(x=500001, color='r', linestyle='--', label='Capping
Point')
plt.xlabel('Median House Value')
plt.ylabel('Frequency')
plt.title('Distribution of Median House Value')
plt.legend()
plt.show()
```

Capped values: 965 out of 20640 (4.68%)



```
housing.shape
```

```
(20640, 9)
```

```
housing=housing.sample(20640)
```

*# If you are curious to know why i do this let me explain, I do this to make data more random so during train test spilt data spread properly into x\_train,x\_test*

```
housing
```

	longitude	latitude	housing_median_age	median_income	\
3478	-118.16	33.88	30.0	2.9779	
18741	-122.67	38.43	17.0	3.2813	
15576	-118.39	34.23	43.0	2.1518	
1398	-121.64	36.66	24.0	5.2285	
1363	-117.00	32.67	16.0	6.6143	
...	...	...	...	...	
1613	-118.33	33.96	42.0	2.3000	
19104	-117.06	32.76	38.0	3.2188	

11074	-120.67	38.76	35.0	2.1682
10334	-119.00	35.39	51.0	2.8295
1244	-122.33	37.55	51.0	9.3694

	median_house_value	ocean_proximity	rooms_per_household	\
3478	169500.0	<1H OCEAN	4.422977	
18741	202700.0	<1H OCEAN	4.980149	
15576	161600.0	<1H OCEAN	3.728125	
1398	248100.0	<1H OCEAN	5.932710	
1363	264100.0	NEAR OCEAN	7.386139	
...	...	...	...	
1613	189200.0	<1H OCEAN	5.285266	
19104	150500.0	NEAR OCEAN	5.571942	
11074	138100.0	INLAND	5.260000	
10334	72100.0	INLAND	4.576667	
1244	500001.0	NEAR OCEAN	8.300971	

	bedrooms_per_room	population_per_household
3478	0.234947	3.083551
18741	0.199302	2.220844
15576	0.250629	3.700000
1398	0.159420	2.740187
1363	0.137176	3.306931
...	...	...
1613	0.214116	2.310345
19104	0.185926	2.287770
11074	0.191540	2.650000
10334	0.206846	2.160000
1244	0.129435	2.815534

[20640 rows x 9 columns]

housing.describe()

	longitude	latitude	housing_median_age	
median_income \				
count	20640.000000	20640.000000	20640.000000	20640.000000
mean	-119.569704	35.631861	28.639486	3.870671
std	2.003532	2.135952	12.585558	1.899822
min	-124.350000	32.540000	1.000000	0.499900
25%	-121.800000	33.930000	18.000000	2.563400
50%	-118.490000	34.260000	29.000000	3.534800
75%	-118.010000	37.710000	37.000000	4.743250
max	-114.310000	41.950000	52.000000	15.000100



	median_house_value	rooms_per_household	bedrooms_per_room \
count	20640.000000	20640.000000	20433.000000
mean	206855.816909	5.429000	0.213039
std	115395.615874	2.474173	0.057983
min	14999.000000	0.846154	0.100000
25%	119600.000000	4.440716	0.175427
50%	179700.000000	5.229129	0.203162
75%	264725.000000	6.052381	0.239821
max	500001.000000	141.909091	1.000000

	population_per_household
count	20640.000000
mean	3.070655
std	10.386050
min	0.692308
25%	2.429741
50%	2.818116
75%	3.282261
max	1243.333333

```

from sklearn.model_selection import train_test_split
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import FunctionTransformer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import LabelBinarizer
from sklearn.pipeline import FeatureUnion

x_train,x_test,y_train,y_test=train_test_split(housing.drop(columns=['
median_house_value']),housing['median_house_value'],test_size=0.20,ran
dom_state=42)

x_train.head()

```

	longitude	latitude	housing_median_age	median_income
ocean_proximity \				
16892 <1H OCEAN	-118.29	33.90	27.0	1.7714
8306 NEAR BAY	-122.31	38.31	32.0	3.8796
5385 NEAR OCEAN	-119.29	34.26	32.0	3.6007
15958 <1H OCEAN	-118.11	33.83	36.0	4.2703
7781	-117.11	32.67	52.0	1.4844

NEAR OCEAN

	rooms_per_household	bedrooms_per_room
population_per_household		
16892	2.532500	0.388944
2.667500		
8306	5.765101	0.177726
2.621924		
5385	5.491667	0.231866
2.240000		
15958	5.966555	0.169843
3.224080		
7781	3.943662	0.253571
3.056338		

x\_test.head()

	longitude	latitude	housing_median_age	median_income
ocean_proximity \				
11890	-117.89	34.07	35.0	3.9808
<1H OCEAN				
14814	-121.95	37.94	21.0	6.8642
INLAND				
11959	-117.40	33.95	32.0	2.4408
INLAND				
4180	-118.08	33.76	27.0	2.0952
<1H OCEAN				
1915	-117.37	34.12	32.0	3.8398
INLAND				

	rooms_per_household	bedrooms_per_room
population_per_household		
11890	5.309963	0.181376
2.966790		
14814	7.315545	0.130352
3.058005		
11959	4.457207	0.248105
2.148649		
4180	3.412903	0.300567
1.245161		
1915	6.230469	0.178056
3.152344		

x\_test['ocean\_proximity'].value\_counts()

ocean_proximity	
<1H OCEAN	1878
INLAND	1296
NEAR OCEAN	514
NEAR BAY	438

```
ISLAND          2
Name: count, dtype: int64
```

```
num_attribute=list(housing.drop(columns=['ocean_proximity','median_hou
se_value']))
```

```
num_attribute
```

```
['longitude',
 'latitude',
 'housing_median_age',
 'median_income',
 'rooms_per_household',
 'bedrooms_per_room',
 'population_per_household']
```

```
cat_attribute=['ocean_proximity']
```

```
pipeline1=Pipeline(steps=[
    ('imputer',SimpleImputer(strategy='median')),
    ('Scaler',StandardScaler())
```

```
],verbose=True)
```

```
pipeline2=Pipeline(steps=[
    ('Label
Encoder',OneHotEncoder(sparse_output=False,drop='first',handle_unknown
='ignore'))
])
```

```
preprocessor=ColumnTransformer(transformers=[
    ('Pipeline 1 For Numerical Columns ',pipeline1,num_attribute),
    ('Encoder For Cateogrical Columns',pipeline2,cat_attribute)
])
```

```
preprocessor
```

```
ColumnTransformer(transformers=[('Pipeline 1 For Numerical Columns ',
                                Pipeline(steps=[('imputer',
```

```
SimpleImputer(strategy='median')),
                                ('Scaler',
```

```
StandardScaler())]),
```

```
                                verbose=True),
                                ['longitude', 'latitude',
```

```
'housing_median_age',
                                'median_income',
```

```

'rooms_per_household',
                                'bedrooms_per_room',
                                'population_per_household']],
                                ('Encoder For Categorical Columns',
                                Pipeline(steps=[('Label Encoder',
OneHotEncoder(drop='first',
handle_unknown='ignore',
sparse_output=False))])),
                                ['ocean_proximity']]))

x_train.shape
(16512, 8)
x_train_transformed=preprocessor.fit_transform(x_train)
[Pipeline] ..... (step 1 of 2) Processing imputer, total= 0.0s
[Pipeline] ..... (step 2 of 2) Processing Scaler, total= 0.0s
x_train_transformed.shape
(16512, 11)
x_test.shape
(4128, 8)
x_test_transformed=preprocessor.fit_transform(x_test)
[Pipeline] ..... (step 1 of 2) Processing imputer, total= 0.0s
[Pipeline] ..... (step 2 of 2) Processing Scaler, total= 0.0s
x_test_transformed.shape
(4128, 11)
x_test['ocean_proximity'].value_counts()
ocean_proximity
<1H OCEAN    1878
INLAND       1296
NEAR OCEAN    514
NEAR BAY      438
ISLAND         2
Name: count, dtype: int64

from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor

```

```

from sklearn.svm import SVR
from sklearn.metrics import
accuracy_score,mean_absolute_error,mean_squared_error,mean_absolute_pe
rcentage_error,root_mean_squared_error
from sklearn.model_selection import cross_val_score
from sklearn.metrics import r2_score, explained_variance_score

```

#### First Model as Linear

##### Regression

```

lr=LinearRegression()
lr.fit(x_train_transformed,y_train)

LinearRegression()

y_pred_lr=lr.predict(x_test_transformed)

print('Percentage
Error',mean_absolute_percentage_error(y_test,y_pred_lr))
print('Mean Square Error',mean_squared_error(y_test,y_pred_lr))
print('RMSE',root_mean_squared_error(y_test,y_pred_lr))
print('R2 Score',r2_score(y_test,y_pred_lr))
print('Variance Score',explained_variance_score(y_test,y_pred_lr))

```

```

Percentage Error 0.3005917558525927
Mean Square Error 5201881955.40867
RMSE 72124.0733417676
R2 Score 0.6094009740883701
Variance Score 0.6094055604744889

```

```

cv_score_lr=cross_val_score(lr,x_train_transformed,y_train,cv=10,scori
ng='neg_mean_squared_error')

```

```

rmse_lr=np.sqrt(-cv_score_lr)
print('CV Scores',rmse_lr)
print('RMSE',rmse_lr.mean())

```

```

CV Scores [69241.34483476 69685.83662162 70012.47076207 69417.97070133
73225.12136675 74577.41502072 70239.71996077 71187.93527516
69418.59489156 70293.43849952]
RMSE 70729.98479342667

```

Conclusion : LinearRegression Model doesn't perform well it goes Underfitting and it fail to learn the pattern cause it give us 27% of mean\_absolute\_percentage\_error and RMSE as 70804 after doing 10 Cross validation and its r2 score is 0.63.

#### Second Model as

##### RandomForestRegressor

```

rfr=RandomForestRegressor(n_estimators=100,max_features=8)
rfr.fit(x_train_transformed,y_train)
RandomForestRegressor(max_features=8)
y_pred_rfr=rfr.predict(x_test_transformed)

print('Percentage
Error',mean_absolute_percentage_error(y_test,y_pred_rfr))
print('Mean Square Error',mean_squared_error(y_test,y_pred_rfr))
print('RMSE',root_mean_squared_error(y_test,y_pred_rfr))
print('R2 Score',r2_score(y_test,y_pred_rfr))
print('Variance Score',explained_variance_score(y_test,y_pred_rfr))

Percentage Error 0.18226256099572993
Mean Square Error 2495179025.639212
RMSE 49951.76699216167
R2 Score 0.8126419428113997
Variance Score 0.8135118348123949

# This line of Code take some time to run

cv_score_rfr=cross_val_score(rfr,x_train_transformed,y_train,cv=10,sco
ring='neg_mean_squared_error')

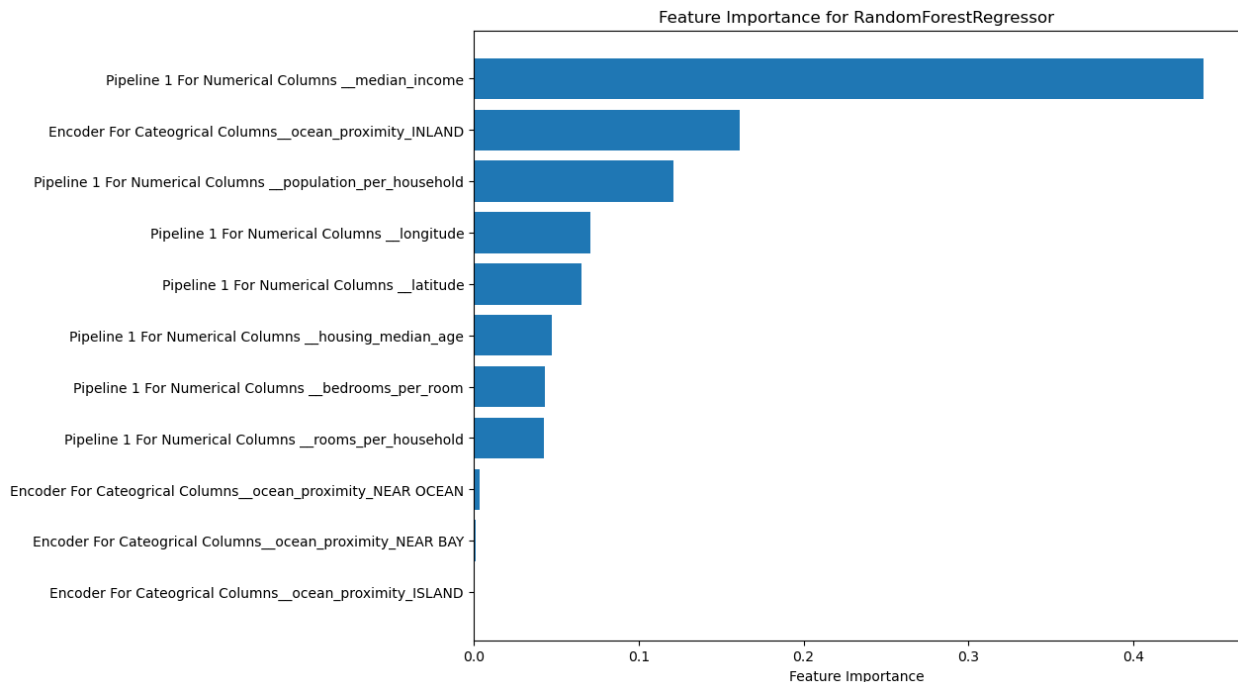
rmse_rfr=np.sqrt(-cv_score_rfr)
print('CV Scores',rmse_rfr)
print('RMSE',rmse_rfr.mean())

CV Scores [46038.38201923 48815.91194992 46998.38371322 49356.48951925
51874.76664592 49718.64121897 50531.29819724 48803.80872544
47660.84540778 47217.64201541]
RMSE 48701.61694123798

importances = rfr.feature_importances_
feature_names = preprocessor.get_feature_names_out()
sorted_indices = importances.argsort()

plt.figure(figsize=(10, 8))
plt.barh(feature_names[sorted_indices], importances[sorted_indices])
plt.xlabel('Feature Importance')
plt.title('Feature Importance for RandomForestRegressor')
plt.show()

```



Random forest Regressor perform well than Linear regression and it give us RMSE of 48094 which is less than of Linear regression. we confirm it by doing cross validation and well R2 score is 0.81 but explained\_variance\_score is 0.81

After doing Cross validation we can confirm that RandomForestRegressor perform better than Linear regressor

### Thrid Model as SVM

```
svr=SVR(kernel='linear')
svr.fit(x_train_transformed,y_train)
SVR(kernel='linear')
y_pred_svr=svr.predict(x_test_transformed)
print('Percentage
Error',mean_absolute_percentage_error(y_test,y_pred_svr))
print('Mean Square Error',mean_squared_error(y_test,y_pred_svr))
print('RMSE',root_mean_squared_error(y_test,y_pred_svr))
print('R2 Score',r2_score(y_test,y_pred_svr))
print('Variance Score',explained_variance_score(y_test,y_pred_svr))
```

```
Percentage Error 0.4935222025503083
Mean Square Error 12384196095.115314
RMSE 111284.30300413133
```



R2 Score 0.07009521305628585  
Variance Score 0.12391550820848884

SVM perform more worst than Linear Regressor

#### Fourth Model as

XGboost

```
from xgboost import XGBRegressor
xgb=XGBRegressor()
xgb.fit(x_train_transformed,y_train)

XGBRegressor(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None,
              early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None,
              feature_types=None,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=None,
              max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=None, max_leaves=None,
              min_child_weight=None, missing=nan,
              monotone_constraints=None,
              multi_strategy=None, n_estimators=None, n_jobs=None,
              num_parallel_tree=None, random_state=None, ...)

xgb_pred=xgb.predict(x_test_transformed)
xgb_pred

array([208091.19, 277352.66, 102060.79, ..., 164733.56, 407450.44,
       336372.7 ], dtype=float32)
```

```
print('Percentage
Error',mean_absolute_percentage_error(y_test,xgb_pred))
print('Mean Square Error',mean_squared_error(y_test,xgb_pred))
print('RMSE',root_mean_squared_error(y_test,xgb_pred))
print('R2 Score',r2_score(y_test,xgb_pred))
print('Variance Score',explained_variance_score(y_test,xgb_pred))
```

Percentage Error 0.17544368021484072  
Mean Square Error 2332399337.424746  
RMSE 48294.92041017095  
R2 Score 0.824864747596245  
Variance Score 0.8251776029971587

```
cv_score_xgb=cross_val_score(xgb,x_train_transformed,y_train,cv=10,sco
ring='neg_mean_squared_error')
```

```
rmse_xgb=np.sqrt(-cv_score_xgb)
print('CV Scores',rmse_xgb)
print('RMSE',rmse_xgb.mean())
```

```
CV Scores [43575.99177336 46011.69672152 45598.91499307 46970.51804133
49974.68446794 46303.37663823 50087.88041028 46612.5573549
44922.44872763 44134.51657066]
RMSE 46419.25856989343
```

Xgboost give us better result than any other algorithms which we use before even it is better than RandomForest Regressor. Xgboost give almost 45943 RMSE with 17% of mean\_absolute\_percentage\_error and its R2 score is about 0.82 with explained Variance score of 0.82 so we choose our final model as XGboost and try to find best parameter for it

```
from sklearn.model_selection import GridSearchCV

parameter=[{'n_estimators':[60,80,100,50], 'max_depth':
[None,5,10,15], 'booster':['gbtree', 'dart']}]

grid_cv=GridSearchCV(estimator=xgb,param_grid=parameter,cv=10,n_jobs=-
1,verbose=3,scoring='neg_mean_squared_error')

grid_cv.fit(x_train_transformed,y_train)

Fitting 10 folds for each of 32 candidates, totalling 320 fits

GridSearchCV(cv=10,
              estimator=XGBRegressor(base_score=None, booster=None,
                                     callbacks=None,
colsample_bylevel=None,
                                     colsample_bynode=None,
                                     colsample_bytree=None,
device=None,
                                     early_stopping_rounds=None,
                                     enable_categorical=False,
eval_metric=None,
                                     feature_types=None, gamma=None,
                                     grow_policy=None,
importance_type=None,
                                     interaction_constraints=None,
                                     learning_rate=None, ...
                                     max_cat_to_onehot=None,
max_delta_step=None,
                                     max_depth=None, max_leaves=None,
                                     min_child_weight=None,
missing=nan,
                                     monotone_constraints=None,
                                     multi_strategy=None,
n_estimators=None,
```

```

n_jobs=None,
num_parallel_tree=None,
random_state=None, ...),
n_jobs=-1,
param_grid=[{'booster': ['gbtree', 'dart'],
              'max_depth': [None, 5, 10, 15],
              'n_estimators': [60, 80, 100, 50]}],
scoring='neg_mean_squared_error', verbose=3)

grid_cv.best_params_
{'booster': 'gbtree', 'max_depth': None, 'n_estimators': 100}

print('Best RMSE Score', np.sqrt(-grid_cv.best_score_))
Best RMSE Score 46465.37840520108


XGB_Model=grid_cv.best_estimator_
XGB_Model
XGBRegressor(base_score=None, booster='gbtree', callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None,
              early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None,
              feature_types=None,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=None,
              max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=None, max_leaves=None,
              min_child_weight=None, missing=nan,
              monotone_constraints=None,
              multi_strategy=None, n_estimators=100, n_jobs=None,
              num_parallel_tree=None, random_state=None, ...)

XGB_Model.feature_importances_
array([0.02714383, 0.02721211, 0.0229763 , 0.21967393, 0.01127589,
       0.01047511, 0.05931564, 0.5555109 , 0.02286633, 0.01276142,
       0.03078858], dtype=float32)

importances = XGB_Model.feature_importances_
feature_names = preprocessor.get_feature_names_out()
sorted_indices = importances.argsort()

plt.figure(figsize=(10, 8))
plt.barh(feature_names[sorted_indices], importances[sorted_indices])
plt.xlabel('Feature Importance')
plt.title('Feature Importance for XGboost')
plt.show()

```



```
XGB_Model.fit(x_train_transformed,y_train)

XGBRegressor(base_score=None, booster='gbtree', callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None,
              early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None,
              feature_types=None,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=None,
              max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=None, max_leaves=None,
              min_child_weight=None, missing=nan,
              monotone_constraints=None,
              multi_strategy=None, n_estimators=100, n_jobs=None,
              num_parallel_tree=None, random_state=None, ...)

new_y_pred_XGB=XGB_Model.predict(x_test_transformed)

print('Percentage
Error',mean_absolute_percentage_error(y_test,new_y_pred_XGB))
print('Mean Square Error',mean_squared_error(y_test,new_y_pred_XGB))
print('RMSE',root_mean_squared_error(y_test,new_y_pred_XGB))
print('R2 Score',r2_score(y_test,new_y_pred_XGB))
print('Variance
Score',explained_variance_score(y_test,new_y_pred_XGB))
```

Percentage Error 0.17544368021484072  
Mean Square Error 2332399337.424746  
RMSE 48294.92041017095  
R2 Score 0.824864747596245  
Variance Score 0.8251776029971587

```
cv_score_xgb_model=cross_val_score(XGB_Model,x_train_transformed,y_train,cv=10,scoring='neg_mean_squared_error')
```

```
rmse_xgb_model=np.sqrt(-cv_score_xgb_model)
print('CV Scores',rmse_xgb_model)
print('RMSE',rmse_xgb_model.mean())
```

```
CV Scores [43575.99177336 46011.69672152 45598.91499307 46970.51804133
 49974.68446794 46303.37663823 50087.88041028 46612.5573549
 44922.44872763 44134.51657066]
RMSE 46419.25856989343
```

After doing GridSearchCV we find out best estimators we are able to get the best result which we can get by reducing mean\_absolute\_percentage\_error to 17% and final RMSE approx to 46419 with r2 score of 0.82 and also explained var which is better than any other model which we used before

```
Final_Model=Pipeline(steps=[
    ('Pre Processing',preprocessor),
    ('Random Forest Regressor',XGB_Model),
])

Final_Model

Pipeline(steps=[('Pre Processing',
                  ColumnTransformer(transformers=[('Pipeline 1 For
Numerical ',
                                                  'Columns ',
Pipeline(steps=[('imputer',
SimpleImputer(strategy='median'))],
('Scaler',
StandardScaler())]),
verbose=True),
['longitude',
'latitude',
'housing_median_age',
'median income'
```

```

'rooms_per_household',
'bedrooms_per_room',
'population_per_household']],
                                ('Encoder For Cat...
                                feature_types=None, gamma=None,
grow_policy=None,
                                importance_type=None,
                                interaction_constraints=None,
learning_rate=None,
                                max_bin=None, max_cat_threshold=None,
                                max_cat_to_onehot=None,
max_delta_step=None,
                                max_depth=None, max_leaves=None,
                                min_child_weight=None, missing=nan,
                                monotone_constraints=None,
multi_strategy=None,
                                n_estimators=100, n_jobs=None,
                                num_parallel_tree=None,
random_state=None, ...)))]
Final_Model.fit(x_train,y_train)
[Pipeline] ..... (step 1 of 2) Processing imputer, total= 0.0s
[Pipeline] ..... (step 2 of 2) Processing Scaler, total= 0.0s
Pipeline(steps=[('Pre Processing',
                  ColumnTransformer(transformers=[('Pipeline 1 For
Numerical ',
                                                  'Columns ',

Pipeline(steps=[('imputer',
SimpleImputer(strategy='median')),
('Scaler',
StandardScaler())]),
verbose=True),
                                ['longitude',
'latitude',
'housing_median_age',
                                'median_income',
'rooms_per_household',
'bedrooms_per_room',

```

```

'population_per_household']],
                                ('Encoder For Cat...
                                feature_types=None, gamma=None,
grow_policy=None,
                                importance_type=None,
                                interaction_constraints=None,
learning_rate=None,
                                max_bin=None, max_cat_threshold=None,
                                max_cat_to_onehot=None,
max_delta_step=None,
                                max_depth=None, max_leaves=None,
                                min_child_weight=None, missing=nan,
                                monotone_constraints=None,
multi_strategy=None,
                                n_estimators=100, n_jobs=None,
                                num_parallel_tree=None,
random_state=None, ...))])

Final_pred_using_xgb=Final_Model.predict(x_test)


print('Final Model RMSE
',root_mean_squared_error(y_test,Final_pred_using_xgb))

Final Model RMSE 45608.15638318089


residuals = y_test - Final_pred_using_xgb
plt.figure(figsize=(8, 6))
sns.scatterplot(x=Final_pred_using_xgb, y=residuals)
plt.axhline(y=0, color='r', linestyle='--')
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
plt.title('Residual Plot')
plt.show()

```





```
plt.figure(figsize=(8, 6))
sns.scatterplot(x=y_test, y=Final_pred_using_xgb)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()],
'r--')
plt.xlabel('True Values')
plt.ylabel('Predicted Values')
plt.title('Prediction Error Plot')
plt.show()
```



```
import pickle

with open('Final Model For California Housing Price','wb') as f:
    pickle.dump(Final_Model,f)

with open('Final Model For California Housing Price','rb') as a:
    model=pickle.load(a)

model.predict(pd.DataFrame(np.array([-117.28,32.74,33.0,2.7515,'NEAR
OCEAN',4.235772,0.266795,1.814024]).reshape(1,8),columns=housing.drop(
columns=['median_house_value']).columns))

array([276127.06], dtype=float32)
```

housing

	longitude	latitude	housing_median_age	median_income	\
3478	-118.16	33.88	30.0	2.9779	
18741	-122.67	38.43	17.0	3.2813	

15576	-118.39	34.23		43.0	2.1518		
1398	-121.64	36.66		24.0	5.2285		
1363	-117.00	32.67		16.0	6.6143		
...	...	...		...	...		
1613	-118.33	33.96		42.0	2.3000		
19104	-117.06	32.76		38.0	3.2188		
11074	-120.67	38.76		35.0	2.1682		
10334	-119.00	35.39		51.0	2.8295		
1244	-122.33	37.55		51.0	9.3694		
	median_house_value	ocean_proximity	rooms_per_household	\			
3478	169500.0	<1H OCEAN	4.422977				
18741	202700.0	<1H OCEAN	4.980149				
15576	161600.0	<1H OCEAN	3.728125				
1398	248100.0	<1H OCEAN	5.932710				
1363	264100.0	NEAR OCEAN	7.386139				
...	...	...	...				
1613	189200.0	<1H OCEAN	5.285266				
19104	150500.0	NEAR OCEAN	5.571942				
11074	138100.0	INLAND	5.260000				
10334	72100.0	INLAND	4.576667				
1244	500001.0	NEAR OCEAN	8.300971				
	bedrooms_per_room	population_per_household					
3478	0.234947	3.083551					
18741	0.199302	2.220844					
15576	0.250629	3.700000					
1398	0.159420	2.740187					
1363	0.137176	3.306931					
...	...	...					
1613	0.214116	2.310345					
19104	0.185926	2.287770					
11074	0.191540	2.650000					
10334	0.206846	2.160000					
1244	0.129435	2.815534					

[20640 rows x 9 columns]